

Robot Navigation in a Simulated Escape Room

Gianluca Viviano Etienne Orio Sara Mashhadi Alizadeh

Abstract—This report describes an autonomous robot system built for a simulated escape-room task. A DJI RoboMaster EP robot, running in CoppeliaSim, must explore an unknown room, find a coloured key cylinder, pick it up, and place it on a pressure plate. The robot uses only wheel-encoder odometry, a 2D lidar scan, and a monocular camera. We built three components on top of ROS 2, Nav2, and `slam_toolbox`: a lidar sensor script injected into the simulator, a colour-based landmark detector using monocular depth estimation, and a nine-state mission controller that combines Nav2 path planning with visual servoing for the pickup. The robot reliably completes all tested scenarios from start to finish.

I. INTRODUCTION

This project is an escape-room task in simulation. A DJI RoboMaster EP robot is placed in a room with obstacles and must find a magenta cylinder (the key), pick it up, and drop it on a green pressure plate to finish. The room layout and object positions are not known in advance.



Fig. 1. The three scenario variants (easy, medium, hard) shown from above in CoppeliaSim. Obstacle count and placement increase across scenarios; the robot spawn, key, and plate positions also change.

The simulation runs in CoppeliaSim with the `robomaster_sim` plugin and `robomaster_ros` driver. We built three custom components on top of `slam_toolbox` and Nav2: a lidar sensor, a colour detector, and a mission controller.

The rest of this report is structured as follows. Section II explains the overall system. Sections III–V describe each component. Section VI covers the problems we ran into and how we fixed them. Section VII concludes.

II. SYSTEM ARCHITECTURE

A. Node graph and data flow

Fig. 2 shows the data flow between all nodes. The system has two custom ROS 2 nodes on top of the standard Nav2 and `slam_toolbox` stacks. `color_detector_node` subscribes to the raw camera image, runs the HSV pipeline, and publishes a latched `PoseStamped` for each detected landmark on `/targets/{cube,plate,door}`; it also streams the live cube bearing on `/targets/cube_live` at every frame for the visual-servo phase. `explorer_node` is the mission FSM: it subscribes to the target poses and the occupancy map, drives long-range navigation by sending `NavigateToPose` goals to Nav2, and takes direct `cmd_vel` control for the short-range pickup and drop manoeuvres.



Fig. 2. Node graph. White boxes: our components. Grey: third-party packages. Arrows show the main ROS 2 topics and actions between nodes.

B. Scenario configuration

The room, obstacles, and task objects are all defined in a JSON file that is passed to the scene builder at startup. The file contains: room size and wall thickness; door openings; a list of obstacles (type, position, size, colour); the key cylinder; and the pressure plate. The scene builder connects to CoppeliaSim, removes any previous objects, and builds the room from scratch. This means changing the scenario only requires swapping the JSON file, no code changes are needed. Three scenarios are included: `easy.json` (5×4 m room,

three obstacles), `medium.json` (same room, four obstacles with a wall-length divider, offset door, and different robot start), and `hard.json` (increased obstacle density and a more constrained layout). The colours in the JSON must match the HSV thresholds in the colour detector.

C. Simulator–ROS bridge

CoppeliaSim exposes a ZMQ-based remote API that lets external programs call simulator functions as Python method calls at runtime, without any shared memory or special ROS bridge. We use it in three places. A scene-building script (`build_scene.py`) uses the API to build the room from the scenario file at startup. The colour detector reads the camera field-of-view angle once at startup to compute the focal length for depth estimation. The mission controller opens and closes the gripper, and toggles the cube’s *detectable* flag to hide it from the lidar while it is being carried, preventing Nav2 from mapping the held cube as a wall directly in front of the robot.

III. PERCEPTION

A. Lidar sensor

CoppeliaSim does not publish laser scans for ROS 2 by default. We wrote a Lua script that is injected into the simulator at startup and attached to the robot. Each cycle (10 Hz) the script rotates a ray sensor through 360 uniformly spaced angles, reads the hit distance, and publishes the result as a laser scan. The robot body is made invisible to the ray sensor so that rays pass through the chassis and only hit walls and obstacles.

B. Colour landmark detection

To find the three landmarks (cube, plate, door), the colour detector processes each camera frame in HSV colour space. For each target it applies a colour mask, finds connected blobs, and keeps the largest one that exceeds a minimum size. The colour ranges used are: cube (magenta, $H \in [140, 170]$), plate (green, $H \in [40, 80]$), door (blue, $H \in [100, 130]$) (OpenCV convention, $H \in [0, 179]$). Magenta was chosen for the cube because it does not appear in the floor or wall textures of the simulation.

1) *Depth from known height (cube and door)*: The cube and door are upright objects with known real heights (0.20 m and 0.50 m). The pinhole camera model gives depth directly from the blob’s pixel height: an object that is H metres tall and appears h_{px} pixels tall is at depth $z = fH/h_{\text{px}}$, where f is the focal length derived from the camera’s field-of-view angle read from the simulator at startup. The blob centroid is then back-projected to that depth in the camera frame and transformed into the map frame via the TF tree. Frames where the blob touches the image edge are skipped because the visible pixel height is cut off and would give a wrong depth.

2) *Floor-ray intersection (pressure plate)*: The plate is flat on the floor, so its pixel height gives no useful depth. Instead, we cast a ray through the centroid pixel, transform the ray into the map frame, and intersect it with the floor plane ($z = 0$) to recover the plate’s (x, y) position. This approach works



Fig. 3. Camera view with HSV blob detection. Coloured overlays show the detected regions for the cube (magenta), pressure plate (green), and door (blue). The centroid of each blob is used to estimate the object’s position.

regardless of the plate’s size and requires no knowledge of the plate dimensions.

IV. NAVIGATION AND MAPPING

A. SLAM

`slam_toolbox` runs in online mapping mode and builds a 2D occupancy map at 5 cm per cell from the lidar scan and odometry. Scan matching aligns each new scan against the current map, correcting the accumulated wheel-slip drift that raw odometry alone cannot account for on the RoboMaster’s mecanum drive. We disabled loop closure because in this small room it caused sudden jumps of 0.3–0.5 m in the `map→odom` transform when the robot revisited a mapped area, which caused Nav2 to abort its current goal mid-path. Without loop closure the drift over one exploration pass is small enough (sub-cell) to complete the task reliably.

B. Nav2

Nav2 handles all long-range path planning and following. We use the NavFn planner (Dijkstra) with a 0.5 m *planner* tolerance (the maximum distance between the computed path endpoint and the requested goal), and the Regulated Pure Pursuit (RPP) controller at a target speed of 0.15 m/s. RPP was preferred over the DWB controller because it has fewer parameters to tune and produces smooth, predictable trajectories at the low speeds dictated by the small room. A separate *goal checker* with a 0.25 m position tolerance determines when the `NavigateToPose` action reports success; this tolerance is large enough for navigation to waypoints but too loose for precision placement (see Section V). The local costmap is a 2×2 m rolling window with 0.35 m inflation around obstacles. The global costmap covers the full map.



Fig. 4. Occupancy grid map built by slam_toolbox during exploration. White cells are free space, black cells are obstacles, and grey is unexplored. The planned path from Nav2 is shown in green.

C. Frontier exploration

To explore the room the robot uses frontier-based exploration. A frontier cell is a free cell (map value 0) with at least one unknown neighbour (map value -1). Adjacent frontier cells are grouped into clusters by BFS; clusters with fewer than 20 cells are discarded as noise. Each tick the robot sends a `NavigateToPose` goal to Nav2 for the closest cluster centre, so the same planner and controller handle both exploration driving and task navigation. Exploration stops once all three landmarks have been detected; at that point the FSM cancels any active Nav2 goal and transitions to **GO_TO_KEY**.

V. MISSION EXECUTION

A. State machine overview

The mission controller (Fig. 5) runs as a timer callback at 4 Hz. It waits at startup until Nav2, the occupancy map, and the robot's position in the map frame are all available, then advances through ten states:

- 1) **EXPLORE**
- 2) **GO_TO_KEY**
- 3) **PICKUP_OPEN**
- 4) **PICKUP_ALIGN**
- 5) **PICKUP_CLOSE**
- 6) **GO_TO_PLATE**
- 7) **DROP_ALIGN**
- 8) **DROP_OPEN**
- 9) **DROP_BACKUP**
- 10) **DONE**

B. Exploration and navigation to key

EXPLORE: the robot sends `NavigateToPose` goals to the nearest frontier cluster centroid (Section IV). Each time a



Fig. 5. Mission state machine. Shaded states use Nav2 for driving; white states use direct velocity commands.

landmark is detected by the colour detector its pose is stored; the state ends as soon as all three landmarks (cube, plate, door) have been seen at least once and their poses are known.



Fig. 6. EXPLORE state: the robot drives toward the nearest frontier while slam_toolbox builds the map. The planned path and current costmap are visible in RViz.

GO_TO_KEY: Nav2 drives the robot to a standoff point 0.9 m in front of the cube, oriented to face it. Stopping 0.9 m away gives the visual-servo phase enough working range before the cylinder base leaves the camera's field of view. If Nav2 reports failure the goal is retried automatically.

C. Pickup sequence

PICKUP_OPEN: commands the gripper to open fully and waits for confirmation before advancing.

PICKUP_ALIGN: the most complex state. The robot uses the live cube bearing from `/targets/cube_live`: a proportional controller on yaw error turns the robot to keep the cube centred in the image while driving forward. At about 0.9 m the base of the 0.20 m cylinder moves below the camera's lower edge, making the pixel-height depth estimate unreliable. The robot freezes the last valid range reading and switches to an *odometry lunge*: it drives straight ahead by $d_{\text{lunge}} = r - d_{\text{engage}} - d_{\text{margin}}$, where $r = 0.9$ m, $d_{\text{engage}} = 0.177$ m is the gripper reach, and $d_{\text{margin}} = 0.03$ m avoids tipping the cylinder. Odometry is used instead of the SLAM transform because short straight-line odometry is smooth and accurate, whereas the `map→odom` transform can jump.

PICKUP_CLOSE: closes the gripper and waits for confirmation. The cube is then hidden from the lidar (`objectspecialproperty_detectable_all = 0`) so that Nav2 does not see the carried object as a wall in front of the robot.



Fig. 7. PICKUP_ALIGN state: the robot approaches the magenta cylinder using visual servoing. The camera view (inset) shows the cube centred in the frame just before the odometry lunge.

D. Drop sequence

GO_TO_PLATE: Nav2 drives the robot to the plate position stored during exploration. The same retry logic as GO_TO_KEY applies. Because the plate is flat on the floor and does not obstruct the lidar, Nav2 can plan and drive directly to it.

Nav2's 0.25 m goal-checker tolerance is too coarse for a 0.3 m plate.

DROP_ALIGN: the controller switches to direct `cmd_vel`. Proportional control on yaw error faces the



Fig. 8. DROP_OPEN state: the cube has been placed on the green pressure plate. The robot is backing away before turning.

robot toward the plate, then drives it forward until the cube's estimated position — $d_{\text{carried}} = 0.207$ m ahead of `base_link` — is within 0.03 m of the plate centre.

DROP_OPEN: the gripper opens and the cube's lidar visibility is restored so the map reflects its new position on the plate.

DROP_BACKUP: the robot reverses 0.25 m in a straight line measured by odometry before allowing any turning. Turning with an open gripper would sweep the cube off the plate.

DONE: the mission controller stops issuing commands and the robot remains stationary.

VI. ENGINEERING CHALLENGES

A. Replacing a custom stack with Nav2 and slam_toolbox

The first version used a script that read all object positions from the simulator directly, published them as perfect odometry, and used a custom Python mapper with an A* planner. This worked but was unreliable: timing problems in the API calls corrupted the pose, and the mapper had no scan matching. We replaced everything with `slam_toolbox` and Nav2, using real wheel-encoder odometry. This was a lot of work to configure (costmaps, controller tuning, recovery behaviours) but gave us a much more reliable system.

B. Lidar: Python node to Lua script

The first lidar was a Python ROS 2 node that asked the simulator for distances to each obstacle one by one. This was slow and depended on knowing the object list in advance. We replaced it with a Lua script running inside the simulator that rotates a native ray sensor and reads distances directly. This is much faster and works for any scene.

C. Colour detector: fixing the coordinate transform

The first version of the colour detector estimated positions in the robot's local frame using a manual rotation based on the starting heading. This broke as soon as the robot turned. We rewrote it to use the full TF tree, transforming each detection from the camera frame to the map frame. This works at any robot pose.

D. Loop closure jumping

With loop closure enabled, `slam_toolbox` sometimes applied a large correction to the map transform when the robot revisited a known area. This moved the robot's estimated position by 0.3–0.5 m instantly, which made Nav2 abort its current goal. We disabled loop closure. The map quality did not noticeably worsen in this small room.

E. Pickup: waypoint to visual servo to odometry lunge

The first pickup strategy drove to a fixed position offset from the cube's known location. The gripper missed often because the position estimate was off by 5–10 cm. We switched to visual servoing, which keeps the cube centred in the camera during the approach and does not depend on map accuracy.

We then found a second problem: below 0.9 m, the cylinder base disappears from the camera's view, making the depth estimate unreliable. We added the odometry lunge to handle this last stretch, where odometry is accurate enough for the short straight-line distance involved.

F. CoppeliaSim Euler angle bug

When setting object orientations using Euler angles, CoppeliaSim silently reorders the angles in some cases (swapping roll and yaw), so the robot spawned facing the wrong direction. We fixed this by using the rotation matrix API instead, which has no ambiguity.

G. Gripper: friction-only hold

The stock RoboMaster gripper script in CoppeliaSim calls `sim.setObjectParent` to reparent any grasped object as a child of the gripper's attach point. With reparenting, the cube's orientation is locked to the gripper: any yaw rotation during driving would tilt the cube and cause it to miss the plate on release. We blocked `setObjectParent` entirely by overriding it with a no-op in a Lua shim that the scene builder injects into the gripper script, and instead rely on friction contact alone to hold the cube. To make this reliable we raised the friction coefficient to 1.0 on the gripper finger shapes and the cube (across all supported physics engines) and added linear and angular damping to the cube so it does not bounce or spin when first squeezed.

H. Removing the GO_TO_DOOR state

We originally included a **GO_TO_DOOR** state that drove the robot through the door after placing the cube on the plate. In testing, the monocular depth estimate for the door (height 0.50 m) had positional errors of 15–20 cm at the approach distances required for navigation. The door opening is only

0.80 m wide, so that error margin made reliable passage through the doorway impossible without repeated replanning. We removed the state entirely: the door is detected during exploration to count as one of the three required landmarks, but the mission ends at **DONE** once the cube is on the plate.

I. CoppeliaSim real-time mode watchdog

The `robomaster_sim` plugin uses a 3-second *simulation-time* heartbeat watchdog. Without real-time mode, CoppeliaSim advances roughly 3× faster than wall clock, so the driver's 1 Hz heartbeat times out in about one real second, silently dropping `/odom`, `/imu`, and all chassis topics for the rest of the run, while the simulator itself continued normally. The failure mode was a robot that stopped moving partway through exploration with no error in the ROS log. The fix was to call `sim.SetBoolParam(boolparam_realtime_simulation, True)` before starting the simulation, which is handled in `run.sh`.

VII. CONCLUSION

We built an autonomous robot system that solves a simulated escape-room task: it explores an unknown room, finds three colour-coded objects, picks up the key cylinder with visual servoing, and drops it on the pressure plate.

The system successfully completed all three scenario variants in testing. Exploration and navigation were reliable across runs; the most sensitive step was the visual-servo pickup, which occasionally struggled when the cube was placed close to a wall and the approach angle was narrow. The drop alignment was more consistent because it uses direct position control rather than camera-based depth estimation.

The main lesson is that using an existing stack (Nav2 and `slam_toolbox`) is much better than building a custom planner, even though configuring it took significant effort. The visual-servo pickup and the odometry lunge were necessary to make the manipulation reliable, because map-based position estimates were not accurate enough close to the object.

Future improvements could include better recovery when the cube is dropped during transport, support for multi-room layouts, or a depth camera to make the pickup more robust.