

Robot Navigation in a Simulated Escape Room

Student 1

*Faculty of Informatics
Università della Svizzera italiana
Lugano, Switzerland*

Student 2

*Faculty of Informatics
Università della Svizzera italiana
Lugano, Switzerland*

Student 3

*Faculty of Informatics
Università della Svizzera italiana
Lugano, Switzerland*

Abstract—This report describes an autonomous robot system built for a simulated escape-room task. A DJI RoboMaster EP robot, running in CoppeliaSim, must explore an unknown room, find a coloured key cylinder, pick it up, and place it on a pressure plate. The robot uses only wheel-encoder odometry, a 2D lidar scan, and a monocular camera. We built three components on top of ROS 2, Nav2, and slam_toolbox: a lidar sensor script injected into the simulator, a colour-based landmark detector using monocular depth estimation, and a nine-state mission controller that combines Nav2 path planning with visual servoing for the pickup. The robot completes the full task from start to finish.

I. INTRODUCTION

This project is an escape-room task in simulation. A DJI RoboMaster EP robot is placed in a room with obstacles and must find a magenta cylinder (the key), pick it up, and drop it on a green pressure plate to finish. The room layout and object positions are not known in advance.



Fig. 1. The three scenario variants (easy, medium, hard) shown from above in CoppeliaSim. Obstacle count and placement increase across scenarios; the robot spawn, key, and plate positions also change.

The simulation runs in CoppeliaSim with the `robomaster_sim` plugin and `robomaster_ros` driver. We built three custom components on top of `slam_toolbox` and Nav2: a lidar sensor, a colour detector, and a mission controller.

The rest of this report is structured as follows. Section II explains the overall system. Sections III–V describe each component. Section VI covers the problems we ran into and how we fixed them. Section VII concludes.

II. SYSTEM ARCHITECTURE

A. Simulator–ROS bridge

CoppeliaSim provides a remote API over ZMQ. We use it in three places: (1) a scene-building script that constructs the room from a JSON file at startup; (2) the colour detector reads the camera field-of-view angle once at startup for the depth calculation; (3) the mission controller opens and closes the gripper and hides or shows the cube to the lidar as needed.

B. Node graph and data flow



Fig. 2. Node graph. White boxes: our components. Grey: third-party packages. Arrows show the main ROS 2 topics and actions between nodes.

Fig. 2 shows the data flow between all nodes.

C. Scenario configuration

The room, obstacles, and task objects are all defined in a JSON file that is passed to the scene builder at startup. The file contains: room size and wall thickness; door openings; a list of obstacles (type, position, size, colour); the key

cylinder; and the pressure plate. The scene builder connects to CoppeliaSim, removes any previous objects, and builds the room from scratch. This means changing the scenario only requires swapping the JSON file, no code changes are needed. Two scenarios are included: `easy.json` (5×4 m room, three obstacles) and `medium.json` (same room, four obstacles with a wall-length divider, offset door, and different robot start). The colours in the JSON must match the HSV thresholds in the colour detector.

D. Coordinate frames

The TF tree is: `map → odom → base_link → laser_link`. `slam_toolbox` provides `map → odom`; `robomaster_ros` provides `odom → base_link`; a static publisher provides `base_link → laser_link` at $z = 0.12$ m. The camera frame `camera_optical_link` is also published by `robomaster_ros` and is used by the colour detector to convert pixel detections to 3D positions.

III. PERCEPTION

A. Lidar sensor

CoppeliaSim does not publish laser scans for ROS 2 by default. We wrote a Lua script that is injected into the simulator at startup and attached to the robot. Each cycle (10 Hz) the script rotates a ray sensor to each of $N = 360$ angles $\theta_i = -\pi + i(2\pi/N)$, reads the hit distance, and publishes the result as a laser scan. The robot body is made invisible to the ray sensor so that rays pass through the chassis and only hit walls and obstacles.

B. Colour landmark detection

To find the three landmarks (cube, plate, door), the colour detector processes each camera frame in HSV colour space. For each target it applies a colour mask, finds connected blobs, and keeps the largest one that exceeds a minimum size. The colour ranges used are: cube (magenta, $H \in [140, 170]$), plate (green, $H \in [40, 80]$), door (blue, $H \in [100, 130]$) (OpenCV convention, $H \in [0, 179]$). Magenta was chosen for the cube because it does not appear in the floor or wall textures of the simulation.

1) *Depth from known height (cube and door)*: The cube and door are upright objects with known real heights (0.20 m and 0.50 m). Using the pinhole camera model, if an object of real height H appears as h_{px} pixels tall in the image, its depth is:

$$z = \frac{f \cdot H}{h_{\text{px}}}, \quad (1)$$

where f is the focal length in pixels, computed from the camera field-of-view angle α : $f = (w/2)/\tan(\alpha/2)$. The centroid pixel (c_x, c_y) is then converted to a 3D point in the camera frame:

$$\mathbf{p}_{\text{cam}} = \left(\frac{c_x - w/2}{f} z, \frac{c_y - h/2}{f} z, z \right), \quad (2)$$



Fig. 3. Camera view with HSV blob detection. Coloured overlays show the detected regions for the cube (magenta), pressure plate (green), and door (blue). The centroid of each blob is used to estimate the object's position.

and transformed to the map frame using the TF tree. Frames where the blob touches the image edge are skipped because the pixel height would be cut off and the depth would be wrong.

2) *Floor-ray intersection (pressure plate)*: The plate is flat on the floor, so measuring its pixel height does not give a useful depth. Instead, we shoot a ray through the centroid pixel, transform both the ray origin and direction to the map frame, and find where the ray crosses the floor plane ($z = 0$):

$$s = \frac{-o_z}{p_z - o_z}, \quad (x, y) = (o_x + s(p_x - o_x), o_y + s(p_y - o_y)), \quad (3)$$

where $\mathbf{o}$ is the camera origin and $\mathbf{p}$ is a point one unit along the ray, both in the map frame. This works regardless of the plate's size.

3) *Live cube position for the pickup*: The cube position in the robot frame (`base_link`) is also published every frame on a separate topic. The pickup controller uses this to always have a fresh direction towards the cube.

IV. NAVIGATION AND MAPPING

A. SLAM

`slam_toolbox` runs in online mode and builds a 2D occupancy map at 5 cm per cell from the lidar scan and odometry. We disabled loop closure because in a small 5×4 m room it caused sudden jumps in the map transform when the robot passed through the same area twice. Those jumps caused Nav2 to abort its current goal. Without loop closure, the map is accurate enough for the task.

B. Nav2

Nav2 handles all long-range path planning and following. We use the NavFn planner (Dijkstra) with a 0.5 m goal



Fig. 4. Occupancy grid map built by slam_toolbox during exploration. White cells are free space, black cells are obstacles, and grey is unexplored. The planned path from Nav2 is shown in green.



Fig. 5. Mission state machine. Shaded states use Nav2 for driving; white states use direct velocity commands.

tolerance and the Regulated Pure Pursuit controller at 0.15 m/s. The local costmap is a 2×2 m rolling window with 0.35 m inflation around obstacles. The global costmap covers the full map.

C. Frontier exploration

To explore the room the robot uses frontier-based exploration. A frontier cell is a free cell (map value 0) that has at least one unknown neighbour (map value -1). Adjacent frontier cells are grouped into clusters using BFS; clusters with fewer than 20 cells are ignored. The robot always drives to the closest cluster centre. Exploration stops once all three landmarks have been seen.

V. MISSION EXECUTION

A. State machine

The mission controller (Fig. 5) runs as a state machine at 4 Hz. It waits at startup until Nav2, the map, and the robot position are all available, then runs through nine states: **EXPLORE** $\rightarrow$ **GO_TO_KEY** $\rightarrow$ **PICKUP_OPEN** $\rightarrow$ **PICKUP_ALIGN** $\rightarrow$ **PICKUP_CLOSE** $\rightarrow$ **GO_TO_PLATE** $\rightarrow$ **DROP_ALIGN** $\rightarrow$ **DROP_OPEN** $\rightarrow$ **DROP_BACKUP** $\rightarrow$ **DONE**.

B. Navigation states

GO_TO_KEY sends a Nav2 goal to a point 0.9 m in front of the cube, facing it. **GO_TO_PLATE** sends a goal to the plate position. If Nav2 fails (for example, a partial plan near a wall), the goal is retried automatically.



Fig. 6. EXPLORE state: the robot drives toward the nearest frontier while slam_toolbox builds the map. The planned path and current costmap are visible in RViz.

C. Pickup (PICKUP_ALIGN)

The pickup works in two steps. First, the robot uses the live cube position from the camera: it turns to face the cube and drives forward while keeping it centred (proportional control on the yaw error).

At about 0.9 m from the cube, the bottom of the 0.20 m cylinder moves below the camera's field of view. At that point the pixel height h_{px} in (1) becomes wrong, so the depth estimate is no longer reliable. The robot therefore switches to an *odometry lunge*: it drives straight ahead by a fixed distance

measured from the odometry, not from SLAM (the SLAM map transform can jump, while odometry is smooth over short distances). The lunge distance is: $d_{\text{lunge}} = r - d_{\text{engage}} - d_{\text{margin}}$, where $r = 0.9$ m is the frozen range, $d_{\text{engage}} = 0.177$ m is the gripper reach, and $d_{\text{margin}} = 0.03$ m avoids tipping the thin cylinder. The gripper then closes.



Fig. 7. PICKUP_ALIGN state: the robot approaches the magenta cylinder using visual servoing. The camera view (inset) shows the cube centred in the frame just before the odometry lunge.

While the cube is being carried, it is hidden from the lidar so that Nav2 does not see it as a wall in front of the robot.

D. Drop (*DROP_ALIGN*, *DROP_OPEN*, *DROP_BACKUP*)



Fig. 8. DROP_OPEN state: the cube has been placed on the green pressure plate. The robot is backing away before turning.

Nav2's position tolerance of 0.25 m is too large for a 0.3 m plate. After arriving near the plate, the robot switches to direct velocity control for the final alignment: it faces the plate and drives until the cube's estimated position (gripper reach ahead of the robot) matches the plate centre, within 0.03 m. The gripper then opens and the cube becomes visible to the lidar again. Finally the robot reverses 0.25 m in a straight line before turning, so that the gripper does not push the cube off the plate.

VI. ENGINEERING CHALLENGES

A. Replacing a custom stack with Nav2 and *slam_toolbox*

The first version used a script that read all object positions from the simulator directly, published them as perfect odometry, and used a custom Python mapper with an A* planner. This worked but was unreliable: timing problems in the API calls corrupted the pose, and the mapper had no scan matching. We replaced everything with *slam_toolbox* and Nav2, using real wheel-encoder odometry. This was a lot of work to configure (costmaps, controller tuning, recovery behaviours) but gave us a much more reliable system.

B. Lidar: Python node to Lua script

The first lidar was a Python ROS 2 node that asked the simulator for distances to each obstacle one by one. This was slow and depended on knowing the object list in advance. We replaced it with a Lua script running inside the simulator that rotates a native ray sensor and reads distances directly. This is much faster and works for any scene.

C. Colour detector: fixing the coordinate transform

The first version of the colour detector estimated positions in the robot's local frame using a manual rotation based on the starting heading. This broke as soon as the robot turned. We rewrote it to use the full TF tree, transforming each detection from the camera frame to the map frame. This works at any robot pose.

D. Loop closure jumping

With loop closure enabled, *slam_toolbox* sometimes applied a large correction to the map transform when the robot revisited a known area. This moved the robot's estimated position by 0.3–0.5 m instantly, which made Nav2 abort its current goal. We disabled loop closure. The map quality did not noticeably worsen in this small room.

E. Pickup: waypoint to visual servo to odometry lunge

The first pickup strategy drove to a fixed position offset from the cube's known location. The gripper missed often because the position estimate was off by 5–10 cm. We switched to visual servoing, which keeps the cube centred in the camera during the approach and does not depend on map accuracy.

We then found a second problem: below 0.9 m, the cylinder base disappears from the camera's view, making the depth estimate unreliable. We added the odometry lunge to handle this last stretch, where odometry is accurate enough for the short straight-line distance involved.

F. CoppeliaSim Euler angle bug

When setting object orientations using Euler angles, CoppeliaSim silently reorders the angles in some cases (swapping roll and yaw), so the robot spawned facing the wrong direction. We fixed this by using the rotation matrix API instead, which has no ambiguity.

G. Too many files in the mission node

As we added states, the mission node grew to 13 files. It became harder to understand and debug than if everything had been in fewer files. We reorganised it down to 5 files. Splitting code into many small files was not worth it for a project of this size.

VII. CONCLUSION

We built an autonomous robot system that solves a simulated escape-room task: it explores an unknown room, finds three colour-coded objects, picks up the key cylinder with visual servoing, and drops it on the pressure plate.

The main lesson is that using an existing stack (Nav2 and slam_toolbox) is much better than building a custom planner, even though configuring it took significant effort. The visual-servo pickup and the odometry lunge were necessary to make the manipulation reliable, because map-based position estimates were not accurate enough close to the object.

Future improvements could include better recovery when the cube is dropped during transport, support for multi-room layouts, or a depth camera to make the pickup more robust.